

Testing Report

Xujie Ma, Yiqi Lei, Xinyi Zhang, Jiaqi Yuan, Yue Liang

11/26/2020

In this notebook, we are presenting 3 algorithms:

1. 14 A3+P5 Doubly Robust Estimation + boosted stumps
2. 21 A6+P5 Regression Adjustment + boosted stumps
3. 15 A4 Regression Estimate

```
library(gbm)
library(dplyr)
```

Data Import

```
high <- read.csv('../data/highDim_dataset.csv')
low <- read.csv('../data/lowDim_dataset.csv')
```

Algo 1: 14 A3+P5 Doubly Robust Estimation + boosted stumps

highDim_dataset

```
# train-test split
n <- nrow(high)
n_train <- round(n*(4/5),0)
train_idx <- sample(1:n,n_train)
train_high <- high[train_idx,]
test_high <- high[-train_idx,]
```

Split treatment and control group, and complete regression for each group.

```
treatment.group.high<-high[high$A==1,-2]
control.group.high<-high[high$A==0,-2]

treatment.model.high<-lm(Y~.,data=treatment.group.high)
control.model.high<-lm(Y~.,data=control.group.high)
```

Estimate $m_1(X)$ and $m_0(X)$ for all entries.

```
X.high<-high[-c(1,2)]

high$m1<-predict(treatment.model.high,X.high)
high$m0<-predict(control.model.high,X.high)
```

Get propensity score for all entries using boosted stumps (Gradient Boosting Machine).

Using grid search to get proper parameters for gbm.

```
# grid search
hyper_grid_high1 <- expand.grid(
  n.trees = c(40,50,60),
  shrinkage = c(.01, .05, .1),
```

```

n.minobsinnode = c(5, 10, 15),
bag.fraction = c(.65, .8, 1),
optimal_trees = 0,           # a place to dump results
min_RMSE = 0                 # a place to dump results
)

# randomize data
random_index <- sample(1:nrow(train_high), nrow(train_high))
random_ames_train <- train_high[random_index, ]

# grid search
for(i in 1:nrow(hyper_grid_high1)) {
  # reproducibility
  set.seed(2020)
  # train model
  gbm.tune <- gbm(
    formula = A~.,
    distribution = "bernoulli",
    data = train_high[-1],
    n.trees = hyper_grid_high1$n.trees[i],
    interaction.depth = 1,
    shrinkage = hyper_grid_high1$shrinkage[i],
    n.minobsinnode = hyper_grid_high1$n.minobsinnode[i],
    bag.fraction = hyper_grid_high1$bag.fraction[i],
    train.fraction = .75
  )

  # add min training error and trees to grid
  hyper_grid_high1$optimal_trees[i] <- which.min(gbm.tune$valid.error)
  hyper_grid_high1$min_RMSE[i] <- sqrt(min(gbm.tune$valid.error))
}

hyper_grid_high1 %>%
  dplyr::arrange(min_RMSE) %>%
  head(10)

```

Apply the parameters with min_RMSE (n.trees=60, shrinkage=0.1, n.minobsinnode=10, bag.fraction=1).

```

set.seed(2020)

tm_highe1 <- system.time(
  boost.high<-gbm(A~., data = train_high[-1],
    distribution = "bernoulli",
    n.trees = 60, # the number of trees
    shrinkage = 0.1, # learning rate
    interaction.depth = 1, # total split
    n.minobsinnode = 10,
    bag.fraction = 1
  )
)

```

Calculate propensity scores for all entries in high.csv

```

tm_highe2 <- system.time(
  high$e <- predict(boost.high, X.high, n.trees = 60, type = 'response')
)

```

```
)
```

Calculate each part in doubly robust estimation and count out the final result.

```
tm_highATE1 <- system.time(  
  {high$p1<-ifelse(high$A==1,(high$Y-high$m1)/high$e,0);  
  high$p2<-ifelse(high$A==0,(high$Y-high$m0)/(1-high$e),0);  
  high$result<-high$m1-high$m0+high$p1-high$p2;  
  ATE.high<-mean(high$result)}  
)
```

```
ATE.high
```

```
## [1] -2.959545
```

```
#alternative function, same result
```

```
#tm_highATE2 <- system.time(  
  # ATE.high<-1/n*(sum((high$A*high$Y-(high$A-high$e)*high$m1)/high$e)  
  # -sum(((1-high$A)*high$Y+(high$A-high$e)*high$m0)/(1-high$e))))  
#ATE.high
```

```
# True ATE:
```

```
true_ATE_high <- -3
```

```
# Comparison:
```

```
true_ATE_high - ATE.high
```

```
## [1] -0.04045505
```

```
time_high<-tm_highel[1]+tm_highe2[1]+tm_highATE1[1]  
cat("Time for training gbm=", tm_highel[1], "s \n")
```

```
## Time for training gbm= 0.421 s
```

```
cat("Time for getting propensity score=", tm_highe2[1], "s \n")
```

```
## Time for getting propensity score= 0.012 s
```

```
cat("Time for calculating ATE=", tm_highATE1[1], "s \n")
```

```
## Time for calculating ATE= 0.001 s
```

lowDim_dataset

```
# train-test split  
n <- nrow(low)  
n_train <- round(n*(4/5),0)  
train_idx <- sample(1:n,n_train)  
train_low <- low[train_idx,]  
test_low <- low[-train_idx,]
```

Split treatment and control group, and complete regression for each group.

```
treatment.group.low<-low[low$A==1,-2]  
control.group.low<-low[low$A==0,-2]  
  
treatment.model.low<-lm(Y~.,data=treatment.group.low)
```

```
control.model.low<-lm(Y~.,data=control.group.low)
```

Estimate $m_1(X)$ and $m_0(X)$ for all entries.

```
X.low<-low[-c(1,2)]
```

```
low$m1<-predict(treatment.model.low,X.low)
```

```
low$m0<-predict(control.model.low,X.low)
```

Get propensity score for all entries using boosted stumps (Gradient Boosting Machine).

Using grid search to get proper parameters for gbm

```
# grid search
hyper_grid_low1 <- expand.grid(
  n.trees = c(40,50,60),
  shrinkage = c(.01, .05, .1),
  n.minobsinnode = c(5, 10, 15),
  bag.fraction = c(.65, .8, 1),
  optimal_trees = 0,          # a place to dump results
  min_RMSE = 0                # a place to dump results
)

# randomize data
random_index <- sample(1:nrow(train_low), nrow(train_low))
random_ames_train <- train_low[random_index, ]

# grid search
for(i in 1:nrow(hyper_grid_low1)) {
  # reproducibility
  set.seed(2020)
  # train model
  gbm.tune <- gbm(
    formula = A~.,
    distribution = "bernoulli",
    data = train_low[-1],
    n.trees = hyper_grid_low1$n.trees[i],
    interaction.depth = 1,
    shrinkage = hyper_grid_low1$shrinkage[i],
    n.minobsinnode = hyper_grid_low1$n.minobsinnode[i],
    bag.fraction = hyper_grid_low1$bag.fraction[i],
    train.fraction = 0.75
  )

  # add min training error and trees to grid
  hyper_grid_low1$optimal_trees[i] <- which.min(gbm.tune$valid.error)
  hyper_grid_low1$min_RMSE[i] <- sqrt(min(gbm.tune$valid.error))
}

hyper_grid_low1 %>%
  dplyr::arrange(min_RMSE) %>%
  head(10)
```

Apply the parameters with min_RMSE (n.trees=60, shrinkage=0.1, n.minobsinnode=15, bag.fraction=0.8).

```

set.seed(2020)

tm_lowe1 <- system.time(
boost.low <- gbm(A~., data = train_low[-1],
  distribution = "bernoulli",
  n.trees = 60, # the number of trees
  shrinkage = 0.1, # learning rate
  interaction.depth = 1, # total split
  n.minobsinnode = 15,
  bag.fraction = 0.8
)
)

```

Calculate propensity scores for all entries in high.csv

```

tm_lowe2 <- system.time(
low$e <- predict(boost.low, X.low, n.trees = 60, type = 'response')
)

```

Calculate each part in doubly robust estimation and count out the final result.

```

tm_lowATE1 <- system.time(
{low$p1<-ifelse(low$A==1,(low$Y-low$m1)/low$e,0);
low$p2<-ifelse(low$A==0,(low$Y-low$m0)/(1-low$e),0);
low$result<-low$m1-low$m0+low$p1-low$p2;
ATE.low<-mean(low$result)}
)

```

```
ATE.low
```

```
## [1] 2.547012
```

```
#alternative function, same result
```

```

#tm_lowATE2 <- system.time(
#ATE.low <- 1/n*(sum((low$A*low$Y-(low$A-low$e)*low$m1)/low$e)
#           -sum(((1-low$A)*low$Y+(low$A-low$e)*low$m0)/(1-low$e))))
#ATE.low

```

```
# True ATE:
```

```
true_ATE_low <- 2.5
```

```
# Comparison:
```

```
true_ATE_low - ATE.low
```

```
## [1] -0.04701199
```

```

time_low<-tm_lowe1[1]+tm_lowe2[1]+tm_lowATE1[1]
cat("Time for training gbm=", tm_lowe1[1], "s \n")

```

```
## Time for training gbm= 0.019 s
```

```
cat("Time for getting propensity score=", tm_lowe2[1], "s \n")
```

```
## Time for getting propensity score= 0.002 s
```

```
cat("Time for calculating ATE=", tm_lowATE1[1], "s \n")
```

```
## Time for calculating ATE= 0.001 s
```

Conclusion for Doubly Robust Estimation

ATE Estimation precision for HighDim dataset and LowDim dataset is pretty similar. Running time for HighDim dataset is around 43 times that for LowDim dataset.

```
table<-data.frame(ATE=c(round(ATE.high,3),round(ATE.low,3)), True.ATE=c(true_ATE_high,true_ATE_low),
                  diff=c(round(true_ATE_high - ATE.high,3),round(true_ATE_low - ATE.low,3)),
                  time=c(time_high,time_low),row.names = c('HighDim', 'LowDim'))

table

##           ATE True.ATE   diff   time
## HighDim -2.960      -3.0 -0.040 0.434
## LowDim   2.547       2.5 -0.047 0.022
```

Algo 2: 21 A6+P5 Regression Adjustment + boosted stumps

The ATEs we got from boosted stumps and regression adjustment are 2.5271 and -3.083. We conclude that it is a close estimate where the true ATEs that are 2.5 and -3.

Methodology and Implementation

1. Reference: D'Agostino RB Jr. Propensity score methods for bias reduction in the comparison of a treatment to a non-randomized control group. Stat Med. 1998 Oct 15;17(19):2265-81. doi: 10.1002/(sici)1097-0258(19981015)17:19<2265::aid-sim918>3.0.co;2-b. PMID: 9802183.
2. Boosted stumps:
 - What is it: an ensemble of weak learners with boosting algorithms. We combine decision tree stumps (decision tree with depth of 1) to predict the propensity score of each sample to simulate a random sample in an observational setting.
 - Implementation: `boost_low = gbm(A~., data = train_low[-1], n.trees = 500, # the number of trees shrinkage = 0.03, # learning rate interaction.depth = 1, # depth of each tree, stumps cv.folds=5)`
3. Regression adjustment:
 - What is it: regress the outcome variable Y on treatment indicator variable A and the estimated propensity score(`pred_high`); We use the estimated coefficient on the A as indicator variable as an estimate of ATE
 - Implementation: `ATE_high = lm(Y~A+pred_high,data=high)`

Data Preparation

```
library(gbm)
library(caret)

## Loading required package: lattice
## Loading required package: ggplot2
high <- read.csv('../data/highDim_dataset.csv')
low <- read.csv('../data/lowDim_dataset.csv')
#high['A'] <- apply(high['A'],1,as.factor)
#low['A'] <- apply(low['A'],1,as.factor)
```

Low Dimension

```
# train-test split
set.seed(2021)
```

```

n <- nrow(low)
n_train <- round(n*(4/5),0)
train_idx <- sample(1:n,n_train)
# test_idx <- setdiff(1:2000, train)
train_low <- low[train_idx,]
test_low <- low[-train_idx,]


## Propensity Score
start0 <- Sys.time()

boost0 = gbm(A~., data = train_low[-1],
  n.trees = 500, # the number of trees, 100, 1000. 10000, no big diff
  shrinkage = 0.03, # learning rate, 0.01, 0.03, 0.05, 0.1
  interaction.depth = 1 # depth of each tree, set 1 as stumps
  ) # here, the parameters we get are from grid search results - see the bottom of the file f

## Distribution not specified, assuming bernoulli ...
# n.trees <- seq(from = 100, to = 10000, by = 100)
# n.trees set the number of trees to be built. Here I choose 1000 manually.
pred0 <- predict(boost0, test_low[-c(1,2)],n.trees = 1000, type = 'response')

## Warning in predict.gbm(boost0, test_low[-c(1, 2)], n.trees = 1000, type =
## "response"): Number of trees not specified or exceeded number fit so far.
## Using 500.
# plot by A to see the distribution of the predicted value
g0_index <- test_low$A == 0
g1_index <- test_low$A == 1
plot(density(pred0[g0_index]),col = 'red')
lines(density(pred0[g1_index]),col = 'blue')
legend('topright',legend = c('group 0','group 1'),fill = c('red','blue'))


```



```
## ATE
```

```
# build a regression model based on the propensity score  
# structure the data frame  
ps0 <- predict(boost0, low[-c(1,2)], n.trees = 100, type = 'response')  
df0 <- data.frame(low$Y, low$A, ps0)  
colnames(df0) <- c('Y', 'A', 'PS')  
model0 <- lm(df0$Y ~ df0$A + df0$PS)  
  
end0 <- Sys.time()
```

```
summary(boost0)
```

```
##      var      rel.inf
## V17 V17 14.6477218
## V22 V22 13.8015173
## V15 V15 12.5012633
## V8  V8  8.4984626
## V12 V12  5.7302940
## V1  V1  5.4095782
## V6  V6  5.0055893
## V3  V3  4.7203187
## V20 V20  4.4765478
## V9  V9  4.2850950
## V2  V2  4.1766691
## V14 V14  4.0724636
## V13 V13  2.6174864
## V5  V5  2.5232861
## V21 V21  2.3838630
## V7  V7  1.4418343
## V4  V4  1.3702618
## V11 V11  0.9535046
## V16 V16  0.8016433
## V19 V19  0.5826000
## V10 V10  0.0000000
## V18 V18  0.0000000
```

```
summary(model0)
```

```
##
## Call:
## lm(formula = df0$Y ~ df0$A + df0$PS)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -10.9058  -2.2172  -0.5321   1.0805  27.7107
##
```

```
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept)  11.5946    0.3589  32.310 < 2e-16 ***
## df0$A        2.5271    0.4101   6.162 1.54e-09 ***
## df0$PS       22.3931    1.4479  15.466 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.495 on 472 degrees of freedom
## Multiple R-squared:  0.4672, Adjusted R-squared:  0.465
## F-statistic: 207 on 2 and 472 DF, p-value: < 2.2e-16
```

```
cat('The total time using boosted stumps and regression adjustment with low dimension data is:', end0 -
```

```
## The total time using boosted stumps and regression adjustment with low dimension data is: 0.1375699
```

High Dimension

```
# train-test split
set.seed(2021)
n <- nrow(high)
n_train <- round(n*(4/5),0)
train_idx <- sample(1:n,n_train)
# test_idx <- setdiff(1:2000, train)
train_high <- high[train_idx,]
test_high <- high[-train_idx,]

## Propensity Score
start1 <- Sys.time()

boost1 = gbm(A~., data = train_high[-1],
             n.trees = 100, # the number of trees
             shrinkage = 0.001, # learning rate
             interaction.depth = 1 # stumps
             ) # here, the parameters we get are from grid search results - see the bottom of the file f
```

```
## Distribution not specified, assuming bernoulli ...
```

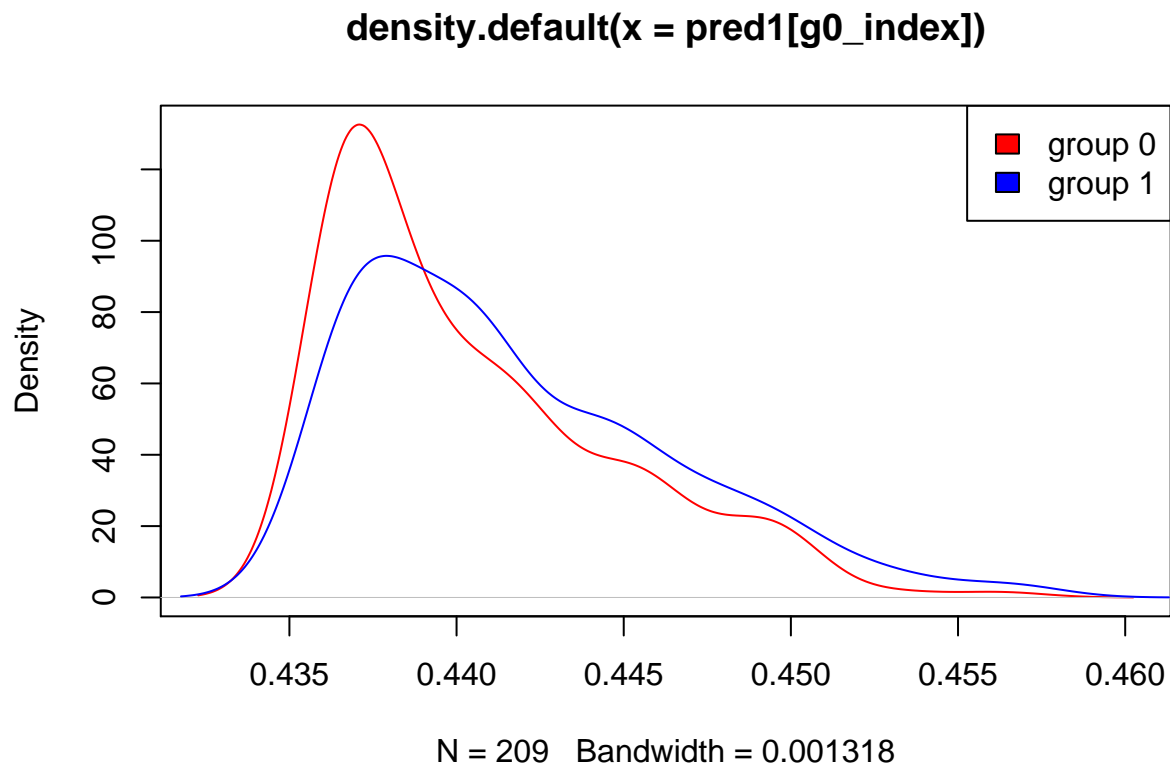
```
#n.trees <- seq(from = 100, to = 10000, by = 100)
# n.trees set the number of trees to be built. Here I choose 1000 manually.
pred1 <- predict(boost1, test_high[-c(1,2)],n.trees = 1000, type = 'response')
```

```
## Warning in predict.gbm(boost1, test_high[-c(1, 2)], n.trees = 1000, type =
## "response"): Number of trees not specified or exceeded number fit so far.
## Using 100.
```

```
length(pred1)
```

```
## [1] 400
```

```
# plot by A to see the distribution of the predicted value
g0_index <- test_high$A == 0
g1_index <- test_high$A == 1
plot(density(pred1[g0_index]),col = 'red')
lines(density(pred1[g1_index]),col = 'blue')
legend('topright',legend = c('group 0','group 1'),fill = c('red','blue'))
```




```
## ATE

# build a regression model based on the propensity score
# structure the data frame
ps1 <- predict(boost1, high[-c(1,2)], n.trees = 100, type = 'response')
df1 <- data.frame(high$Y, high$A, ps1)
colnames(df1) <- c('Y', 'A', 'PS')
model1 <- lm(df1$Y ~ df1$A + df1$PS)

end1 <- Sys.time()

summary(boost1)
```



##	var	rel.inf
##	V95	31.7118710
##	V91	16.5506737
##	V83	9.6377250
##	V92	8.9060494
##	V65	6.5617795
##	V76	5.4475166
##	V63	3.3178531
##	V121	2.9209869
##	V68	1.9309898
##	V67	1.8390144
##	V134	1.7500140
##	V124	1.5023956
##	V73	1.0652995
##	V180	1.0160785
##	V99	0.8926300
##	V90	0.8856194
##	V94	0.8699262
##	V69	0.8450043
##	V131	0.8214455
##	V141	0.7930416
##	V89	0.7340859
##	V1	0.0000000
##	V2	0.0000000
##	V3	0.0000000
##	V4	0.0000000
##	V5	0.0000000
##	V6	0.0000000
##	V7	0.0000000
##	V8	0.0000000
##	V9	0.0000000
##	V10	0.0000000

##	V11	V11	0.0000000
##	V12	V12	0.0000000
##	V13	V13	0.0000000
##	V14	V14	0.0000000
##	V15	V15	0.0000000
##	V16	V16	0.0000000
##	V17	V17	0.0000000
##	V18	V18	0.0000000
##	V19	V19	0.0000000
##	V20	V20	0.0000000
##	V21	V21	0.0000000
##	V22	V22	0.0000000
##	V23	V23	0.0000000
##	V24	V24	0.0000000
##	V25	V25	0.0000000
##	V26	V26	0.0000000
##	V27	V27	0.0000000
##	V28	V28	0.0000000
##	V29	V29	0.0000000
##	V30	V30	0.0000000
##	V31	V31	0.0000000
##	V32	V32	0.0000000
##	V33	V33	0.0000000
##	V34	V34	0.0000000
##	V35	V35	0.0000000
##	V36	V36	0.0000000
##	V37	V37	0.0000000
##	V38	V38	0.0000000
##	V39	V39	0.0000000
##	V40	V40	0.0000000
##	V41	V41	0.0000000
##	V42	V42	0.0000000
##	V43	V43	0.0000000
##	V44	V44	0.0000000
##	V45	V45	0.0000000
##	V46	V46	0.0000000
##	V47	V47	0.0000000
##	V48	V48	0.0000000
##	V49	V49	0.0000000
##	V50	V50	0.0000000
##	V51	V51	0.0000000
##	V52	V52	0.0000000
##	V53	V53	0.0000000
##	V54	V54	0.0000000
##	V55	V55	0.0000000
##	V56	V56	0.0000000
##	V57	V57	0.0000000
##	V58	V58	0.0000000
##	V59	V59	0.0000000
##	V60	V60	0.0000000
##	V61	V61	0.0000000
##	V62	V62	0.0000000
##	V64	V64	0.0000000
##	V66	V66	0.0000000

##	V70	V70	0.0000000
##	V71	V71	0.0000000
##	V72	V72	0.0000000
##	V74	V74	0.0000000
##	V75	V75	0.0000000
##	V77	V77	0.0000000
##	V78	V78	0.0000000
##	V79	V79	0.0000000
##	V80	V80	0.0000000
##	V81	V81	0.0000000
##	V82	V82	0.0000000
##	V84	V84	0.0000000
##	V85	V85	0.0000000
##	V86	V86	0.0000000
##	V87	V87	0.0000000
##	V88	V88	0.0000000
##	V93	V93	0.0000000
##	V96	V96	0.0000000
##	V97	V97	0.0000000
##	V98	V98	0.0000000
##	V100	V100	0.0000000
##	V101	V101	0.0000000
##	V102	V102	0.0000000
##	V103	V103	0.0000000
##	V104	V104	0.0000000
##	V105	V105	0.0000000
##	V106	V106	0.0000000
##	V107	V107	0.0000000
##	V108	V108	0.0000000
##	V109	V109	0.0000000
##	V110	V110	0.0000000
##	V111	V111	0.0000000
##	V112	V112	0.0000000
##	V113	V113	0.0000000
##	V114	V114	0.0000000
##	V115	V115	0.0000000
##	V116	V116	0.0000000
##	V117	V117	0.0000000
##	V118	V118	0.0000000
##	V119	V119	0.0000000
##	V120	V120	0.0000000
##	V122	V122	0.0000000
##	V123	V123	0.0000000
##	V125	V125	0.0000000
##	V126	V126	0.0000000
##	V127	V127	0.0000000
##	V128	V128	0.0000000
##	V129	V129	0.0000000
##	V130	V130	0.0000000
##	V132	V132	0.0000000
##	V133	V133	0.0000000
##	V135	V135	0.0000000
##	V136	V136	0.0000000
##	V137	V137	0.0000000

```
## V138 V138 0.0000000
## V139 V139 0.0000000
## V140 V140 0.0000000
## V142 V142 0.0000000
## V143 V143 0.0000000
## V144 V144 0.0000000
## V145 V145 0.0000000
## V146 V146 0.0000000
## V147 V147 0.0000000
## V148 V148 0.0000000
## V149 V149 0.0000000
## V150 V150 0.0000000
## V151 V151 0.0000000
## V152 V152 0.0000000
## V153 V153 0.0000000
## V154 V154 0.0000000
## V155 V155 0.0000000
## V156 V156 0.0000000
## V157 V157 0.0000000
## V158 V158 0.0000000
## V159 V159 0.0000000
## V160 V160 0.0000000
## V161 V161 0.0000000
## V162 V162 0.0000000
## V163 V163 0.0000000
## V164 V164 0.0000000
## V165 V165 0.0000000
## V166 V166 0.0000000
## V167 V167 0.0000000
## V168 V168 0.0000000
## V169 V169 0.0000000
## V170 V170 0.0000000
## V171 V171 0.0000000
## V172 V172 0.0000000
## V173 V173 0.0000000
## V174 V174 0.0000000
## V175 V175 0.0000000
## V176 V176 0.0000000
## V177 V177 0.0000000
## V178 V178 0.0000000
## V179 V179 0.0000000
## V181 V181 0.0000000
## V182 V182 0.0000000
## V183 V183 0.0000000
## V184 V184 0.0000000
## V185 V185 0.0000000
```

```
summary(model1)
```

```
##
## Call:
## lm(formula = df1$Y ~ df1$A + df1$PS)
##
## Residuals:
```

##	Min	1Q	Median	3Q	Max
----	-----	----	--------	----	-----

```
## -13.4431 -2.6174 0.0114 2.3630 19.8397
##
## Coefficients:
##           Estimate Std. Error t value Pr(>|t|)
## (Intercept) -521.3317      9.6763  -53.88  <2e-16 ***
## df1$A        -3.0830      0.1987  -15.52  <2e-16 ***
## df1$PS       1157.6456     21.9805   52.67  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4.296 on 1997 degrees of freedom
## Multiple R-squared:  0.5824, Adjusted R-squared:  0.5819
## F-statistic: 1392 on 2 and 1997 DF,  p-value: < 2.2e-16
cat('The total time using boosted stumps and regression adjustment with high dimension data is:', end1)
```

```
## The total time using boosted stumps and regression adjustment with high dimension data is: 0.673718
```

The summary of the model gives a feature importance plot. Conduct prediction on the test set so we can have Test Error as an evaluation. The density plot shows the overlap of propensity score between the two groups.

Additional Code: Grid Search

In this part, we included our code for conducting grid search for lowdim. The procedure is similar in highdim. For the readability of the file, we choose to comment out the code.

```
# grid search
# hyper_grid_low1 <- expand.grid(
#   shrinkage = c(.01, 0.03, 0.05),
#   interaction.depth = 1 - since it is boosted stumps
#   n.minobsinnode = c(5, 10, 15),
#   bag.fraction = c(.65, .8, 1),
#   optimal_trees = 0,                # a place to dump results
#   min_RMSE = 0                      # a place to dump results
#)

# randomize data
# random_index <- sample(1:nrow(train_low), nrow(train_low))
# random_ames_train <- train_low[random_index, ]

# grid search
# for(i in 1:nrow(hyper_grid_low1)) {
#   # reproducibility
#   set.seed(2020)
#   # train model
#   gbm.tune <- gbm(
#     formula = A~.,
#     data = train_low[-1],
#     n.trees = 500,
#     interaction.depth = hyper_grid_low1$interaction.depth[i],
#     shrinkage = hyper_grid_low1$shrinkage[i],
#     n.minobsinnode = hyper_grid_low1$n.minobsinnode[i],
#     bag.fraction = hyper_grid_low1$bag.fraction[i],
#     train.fraction = .75,
#     n.cores = NULL, # will use all cores by default
#     verbose = FALSE
```



```

# )

# add min training error and trees to grid
# hyper_grid_low1$optimal_trees[i] <- which.min(gbm.tune$valid.error)
# hyper_grid_low1$min_RMSE[i] <- sqrt(min(gbm.tune$valid.error))
# }

# hyper_grid_low1 %>%
# dplyr::arrange(min_RMSE) %>%
# head(10)

```

Algo 3: 15 A4 Regression Estimate

Understanding

Regression Estimate is a really simple estimation model to calculate ATE, which do not require Propensity Scores calculation. This makes it a straight forward model and a computational efficient model. By implementing the linear regression on treated groups and untreated groups, we could regress on different groups to get the two different sets of parameters and then by predicting the models on the whole dataset, subtracting the prediction we can get the difference between the two regression models. In the end, we can calculate the ATE(Average Treatment Effect) by taking the average of the difference.

$$ATE = N^{-1} \sum_{i=1}^N (\hat{m}_1(X_i) - \hat{m}_0(X_i))$$

Denote that

N is the number of samples in the dataset,

X_i is the datapoint in the dataset,

m_1 is the regression model learned from the treated groups,

m_0 is the regression model learned from the untreated groups,

$\hat{m}_1(X_i)$ is the prediction of the regression model m_1 on the datapoint X_i ,

$\hat{m}_0(X_i)$ is the prediction of the regression model m_0 on the datapoint X_i .

Implementation

Read the data and split the data into two groups – Treated Group and Untreated Group

```

high_data <- read.csv('../data/highDim_dataset.csv')
low_data <- read.csv('../data/lowDim_dataset.csv')

N_high <- dim(high_data)[1]
N_low <- dim(low_data)[1]

high_data_X <- high_data[,3:dim(high_data)[2]]
low_data_X <- low_data[,3:dim(low_data)[2]]

high_treated <- high_data[high_data$A==1,-2]
high_untreated <- high_data[high_data$A==0,-2]

```

```

N_high_treated <- dim(high_treated)[1]
N_high_untreated <- dim(high_untreated)[1]

low_treated <- low_data[low_data$A==1,-2]
low_untreated <- low_data[low_data$A==0,-2]

N_low_treated <- dim(low_treated)[1]
N_low_untreated <- dim(low_untreated)[1]

```

Train the data and record the training time of two datasets

```

time<- system.time({
  high_treated_lm <- lm(Y~.,data = high_treated);
  high_untreated_lm <- lm(Y~.,data = high_untreated);
  high_treated_predict_all <- predict(high_treated_lm,newdata = high_data_X);
  high_untreated_predict_all <- predict(high_untreated_lm,newdata = high_data_X)})
train_time_high <- time[1]
train_time_high

```

```

## user.self
##      0.133

```

```

time<- system.time({
  low_treated_lm <- lm(Y~.,data = low_treated);
  low_untreated_lm <- lm(Y~.,data = low_untreated);
  low_treated_predict_all <- predict(low_treated_lm,newdata = low_data_X);
  low_untreated_predict_all <- predict(low_untreated_lm,newdata = low_data_X)})
train_time_low <- time[1]
train_time_low

```

```

## user.self
##      0.008

```

Calculate the ATE

```

reg_est_ATE_high<-sum(high_treated_predict_all - high_untreated_predict_all)/N_high
reg_est_ATE_low<-sum(low_treated_predict_all - low_untreated_predict_all)/N_low
reg_est_ATE_high

```

```

## [1] -2.95978

```

```

reg_est_ATE_low

```

```

## [1] 2.526944

```

Compare the ATE with the true ATE

```

# True ATE:
true_ATE_high <- -3
true_ATE_low <- 2.5

# Comparison:
abs(true_ATE_high - reg_est_ATE_high) /abs(true_ATE_high)

```

```

## [1] 0.01340679

```

```
abs(true_ATE_low - reg_est_ATE_low) / abs(true_ATE_low)
```

```
## [1] 0.01077759
```

Conclustions

Comparision between the two dataset

We can conclude that the model is more fit to the low dimension dataset. With higher dimension, the ATE has higher bias rate(1.34% vs 1.08%).

Comparision among the three models

The table shows the result of the three algorithm's ATE in the two different datasets.

Algorithm	High ATE	Low ATE	High Train Time	Low Train Time
True ATE	-3	2.5	-	-
Doubly Robust Estimation + Boosted Stumps	-2.9626	2.5187	1.2180	0.0230
Regression Estimate	-2.9598	2.5269	0.2270	0.0190
Regression Adjustment + Boosted Stumps	-3.0830	2.5271	0.5060	0.1287

From the table above, we can clearly conclude that the Regression Estimate's accuracy is relatively high, but slightly lower than the Doubly Robust Estimation + Boosted Stumps model. However, the training time of Doubly Robust Estimation + Boosted Stumps model is higher than the Regression Estimate model for both high dimension dataset and low dimension dataset. We can conclude that the Regression Estimate is more computational efficient but slightly less accuracy than the Doubly Robust Estimation + Boosted Stumps model.